
COMPONENT WORKLOAD EMULATOR UTILIZATION TEST SUITE

User and Developer's Guide

Version 1.0

© 2009 – 2011 by James H. Hill; all rights reserved

Contents

| | |
|---|-----------|
| Preface | v |
| I Getting Started | 1 |
| 1 Building and Installing CUTS | 3 |
| 1.1 CUTS Runtime Toolset | 3 |
| 1.1.1 Obtaining Source Files | 4 |
| 1.1.2 System Configuration | 5 |
| 1.1.3 Installation | 5 |
| 1.2 CUTS Modeling Toolset (Windows-only) | 5 |
| 1.2.1 Installing Prebuilt Version | 6 |
| 1.2.2 Building from Sources (advanced) | 6 |
| 1.3 CUTS Analysis Toolset | 7 |
| 2 Quick Start Tutorial | 9 |
| 2.1 Modeling Behavior and Workload | 9 |
| 2.2 Generating Test System Implementation | 11 |
| 2.3 Executing in Target Environment | 11 |
| 2.4 Analyzing Test Results | 14 |
| II CUTS Modeling Toolset | 15 |
| III CUTS Runtime Toolset | 17 |
| 3 CUTS Node Management Facilities | 19 |
| 3.1 Overview | 19 |

| | | |
|-----------|--|-----------|
| 3.2 | The Node Daemon | 20 |
| 3.2.1 | Configuring the Node Daemon | 20 |
| 3.2.2 | Running the Node Daemon | 22 |
| 3.2.3 | Configuring the Node Daemon Server | 22 |
| 3.2.4 | Node Daemon Schema Definition | 23 |
| 3.3 | The Node Daemon Client | 25 |
| 3.3.1 | Running the Node Daemon Client | 25 |
| 4 | CUTS Testing Facilities | 27 |
| 5 | CUTS Logging Facilities | 29 |
| 5.1 | Overview | 29 |
| 5.2 | The Logging Server | 31 |
| 5.2.1 | Running the Logging Server | 31 |
| 5.2.2 | Configuring the Logging Server | 31 |
| 5.3 | The Logging Client | 32 |
| 5.3.1 | Running the Logging Client | 32 |
| 5.3.2 | Configuring the Logging Client | 32 |
| 5.4 | The Client Loggers | 33 |
| 5.4.1 | Direct Integration | 33 |
| 5.4.2 | Indirect Integration | 35 |
| IV | CUTS Analysis Toolset | 37 |
| V | Appendix | 39 |
| A | Building and Installing Third-Party Libraries | 41 |
| A.1 | Boost | 41 |
| A.1.1 | System Configuration | 41 |
| A.1.2 | Installation | 42 |
| A.2 | Xerces-C | 42 |
| A.2.1 | System Configuration | 42 |
| A.2.2 | Installation | 43 |
| A.3 | DOC Group Middleware | 43 |
| A.3.1 | System Configuration | 43 |
| A.3.2 | Installation | 44 |
| A.3.3 | ACE DataBase Connector (ADBC) Framework | 45 |

| | | |
|----------|--|-----------|
| A.4 | XML Schema Compiler (XSC) | 45 |
| A.4.1 | System Configuration | 45 |
| A.4.2 | Installation | 45 |
| A.5 | Perl-Compatible Regular Expressions (PCRE) | 46 |
| A.5.1 | System Configuration | 46 |
| A.5.2 | Installation | 46 |
| A.6 | SQLite | 46 |
| A.6.1 | System Configuration | 46 |
| A.6.2 | Installation | 47 |
| A.7 | RTI-DDS | 47 |
| A.7.1 | System Configuration | 47 |
| A.7.2 | Installation | 47 |
| A.8 | OpenSplice DDS | 48 |
| A.8.1 | System Configuration | 48 |
| A.8.2 | Installation | 48 |
| B | Autobuilds of CUTS | 51 |
| B.1 | Hosting a Windows Build | 51 |

Preface

Enterprise distributed systems (such as air traffic management systems, cloud computing centers, and shipboard computing environments) are steadily increasing in size (e.g., lines of source code and number of hosts in the target environment) and complexity (e.g., application scenarios). To address challenges associated with developing next-generation enterprise distributed systems, the level-of-abstraction for software development is steadily increasing. Now-a-days, distributed system developers focus more on the system’s “business-logic” instead of wrestling with low-level implementation details, such as development and configuration, resource management, and fault tolerance. Moreover, increasing the level-of-abstraction for software development promotes reuse of the system’s “business-logic” across different application domains, which inherently reduces (re)invention of core intellectual property.

Although increasing the level-of-abstraction for software development is improving functional properties of next-generation enterprise distributed systems, system quality-of-service (QoS) properties (e.g., latency, throughput, and scalability) are not validated until late in the software lifecycle, *i.e.*, at system integration time. This is due in part to the *serialized-phasing development problem* where the infrastructure- and application-level system entities, are developed during different phases of the software lifecycle. Consequently, distributed system developers do not realize the system under development does not meet its QoS requirements until its too late, *i.e.*, at complete system integration time, at the expense of overrun project cost and deadlines.

System execution modeling (SEM) is a model-driven engineering technique that helps overcome the effects of serialized-phasing development. SEM tools provide distributed system developers with the necessary artifacts for modeling system behavior and workload, such as computational attributes, resource requirements, and network communication. The constructed models are then used to validate QoS properties of the system under development during early phases of the software lifecycle. This enables distributed system developers to pinpoint potential QoS bottlenecks before they become too costly to locate and resolve in a cost-effective manner late in the software lifecycle.

The Component Workload Emulator (CoWorkEr) Utilization Test Suite (CUTS) is a SEM tool designed for next-generation enterprise distributed systems. Distributed system developers and testers use CUTS to model the expected behavior and workload of system components under devel-

opment using high-level domain-specific modeling languages. Model interpreters then transform the constructed behavior and workload models into source code for the target architecture, *i.e.*, the auto-generated components same interfaces and attributes as their real counterparts. Finally, system developers and testers emulate the auto-generated components in the target (or representative) environment and collect and analyze QoS metrics. This enables distributed system developers and testers to conduct system integration test at early stages of software lifecycle, instead of waiting until complete system integration time to perform such testing.

This book therefore serves as the user's guide to CUTS. It details its main functionalities, and includes concrete examples to provide better understanding of its concepts. This book is organized organized as follows:

- **Chapter 1** details how to build and install CUTS and its different toolsets. Depending on your needs, you may or may not install all the toolset. It is not a requirement to install all the toolsets.
- **Chapter 2** provides a Quick Start Tutorial on using CUTS. The tutorial uses a real example that goes over the key task in CUTS, such as modeling behavior and workload, generating source code, running experiments, and analyzing collected QoS metrics.

Hope that you enjoy using CUTS as it helps you understand QoS properties during early phases of the software lifecycle.

Part I

Getting Started

Chapter 1

Building and Installing CUTS

CUTS has many small projects that comprise the entire system execution modeling (SEM) tool. Its many projects, however, can be divided into three main toolsets:

- Runtime architecture
- Modeling tools
- Analysis tools

This chapter discusses how to build and install each of the aforementioned toolsets for CUTS. Depending on your usage of CUTS, you may not need to build and install all projects in a category on the same machine. For example, you may install the runtime toolset of CUTS in your testing environment, and the modeling and analysis toolset outside of your testing environment to minimize interference with the testing process when viewing collected performance metrics in real-time. We are aware of these needs, and have setup the build process to decouple projects within a toolset from projects external to its corresponding toolset. You can, therefore, refer to each of the following sections on building and installation in isolation since toolsets do not explicitly depend on each other.

1.1 CUTS Runtime Toolset

The runtime toolset for CUTS allows developers and testers to emulate system experiments on their target architecture. CUTS also provides the mechanisms to monitor and collect performance metrics for the executing system. In order to build the CUTS runtime architecture, you will need the following technologies installed on the target machine(s):

- **DOC Group Middleware** - This set of middleware is used to abstract away complexities associated with implementing applications that operate on many different operating systems

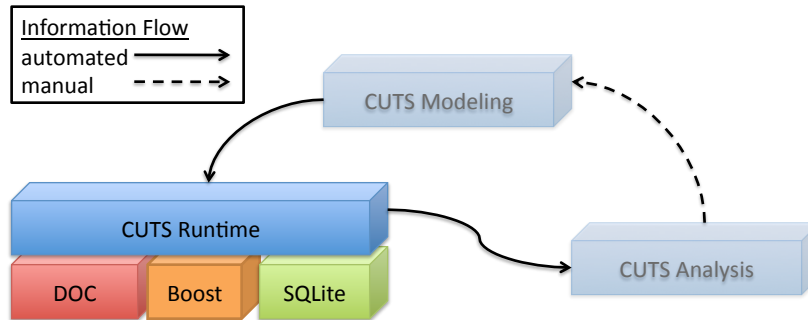


Figure 1.1. Building blocks for the CUTS runtime architecture

(ACE) and perform distributed communication (TAO). You can download DOC Group Middleware from the following location: <http://www.dre.vanderbilt.edu>.

- **Boost** - This set of libraries is used for implementing the parsers (Boost Spirit) used within different artifacts of CUTS. You can download Boost from the following location: <http://www.boost.org>.
- **SQLite** - This set of libraries is used to support flat file archives for test results. You can download SQLite at the following location: <http://www.sqlite.org>.
- **PCRE (not pictured)** - This set of libraries is used to enable PERL regular expression support in CUTS. You can download PCRE at the following location: <http://www.pcre.org>.
- **XSC (not pictured)** - This set of libraries and applications is used to convert XML documents to/from objects. You can download XSC from the following location: <svn://svn.dre.vanderbilt.edu/XSC/trunk>.

1.1.1 Obtaining Source Files

You can obtain the latest snapshot of the CUTS runtime architecture from the DOC Group Subversion repository at the following location:

<svn://svn.dre.vanderbilt.edu/DOC/CUTS/trunk/CUTS>

If you need to download a stable version of the source code, then you can access at one of the subdirectories in the repository at the following location:

<svn://svn.dre.vanderbilt.edu/DOC/CUTS/tags>

1.1.2 System Configuration

Before you can build CUTS, you must for configure you environment. Please set the following environment variables:

```
%> export CUTS_ROOT=location of CUTS
%> export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$CUTS_ROOT/lib
%> export PATH=$PATH:$CUTS_ROOT/bin
```

Please see Appendix A for instructions for configuring, building, and installing third-party libraries needed by CUTS, such as DOC Group Middleware, Boost, and SQLite.

1.1.3 Installation

We use Makefile, Workspace, Project Creator (MPC) ¹ to assist in building CUTS on different operating systems, and with different compilers. Before you can build the runtime architecture, you must first generate the target workspace. Use the following command to generate the workspace:

```
%> $ACE_ROOT/bin/mwc.pl -type TYPE [-features FEATURES] CUTS.mwc
```

where TYPE is your compiler type, and FEATURES is a comma-separated list of features for your build of the CUTS runtime architecture.² The complete set of features for CUTS is located in \$CUTS_ROOT/default.features.tmpl. It is recommended that you copy the template feature file to \$CUTS_ROOT/default.features and set the appropriate features for your workspace by modifying the new file. This way you do not have to use the -features command-line option unless you want to override your default feature selection. Once you have generated the workspace, you can build the solution using your specified compiler.

1.2 CUTS Modeling Toolset (Windows-only)

CUTS modeling tools provide system developers and testers with an enviroment for rapidly constructing experiments for distributed component-based systems, and generating testing system for their target architecture. The modeling tools are built on the following technologies:

- **GME** - The Generic Modeling Environment (GME) is a graphical modeling tool for creating domain-specific modeling languages (DSMLs). You can download GME from the following location: <http://www.isis.vanderbilt.edu/projects/GME>.

¹For more information on MPC, please see the following location: <http://www.ociweb.com/products/mpc>

²You can type \$ACE_ROOT/bin/mwc.pl --help to view the command-line options for MPC.

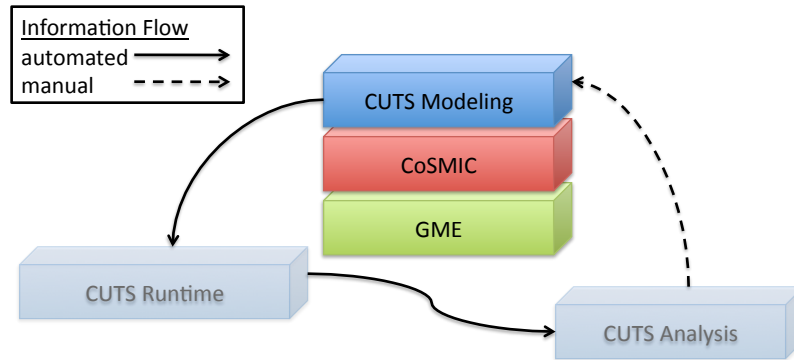


Figure 1.2. Building blocks for the CUTS modeling tools

- **CoSMIC** - The Component Synthesis via Model Integrated Computing (CoSMIC) is a tool suite designed to address the complexities of building large-scale component-based systems, such as system assembly, deployment, packaging, and planning. You can download CoSMIC from the following location: <http://www.dre.vanderbilt.edu/cosmic>.

1.2.1 Installing Prebuilt Version

The easiest and quickest way to install the CUTS modeling tools is to install them using the .msi installer. You can download the latest version of the CUTS modeling tools from the following location: <http://www.dre.vanderbilt.edu/CUTS/downloads>. Before you can install the CUTS modeling tools, please make sure you have installed GME and CoSMIC. Once both GME and CoSMIC are installed (in that order), then you can install the CUTS modeling tools.

1.2.2 Building from Sources (advanced)

If you choose, you can build the modeling tools from source. This is given you have downloaded all the source from the CUTS source code repository (see Section 1.1.1). To build the CUTS modeling tools, first install the latest version of GME. Then, you **MUST** build CoSMIC from source as well. This is required because the CUTS modeling tools are built on top of CoSMIC, and its therefore dependent on CoSMIC. You can find instructions for building CoSMIC from source at the following location:

<http://www.dre.vanderbilt.edu/cosmic/downloads>

After you have built and installed CoSMIC from sources, you are ready to build the CUTS modeling tools from source. Please use the following commands to build the modeling tools using your flavor of Visual C++:

```
%> cd %CUTS_ROOT%  
%> mwc.pl -type [vc type] -features modeling=1,cosmic=1 CUTS_CoSMIC.mwc  
%> open solution and build
```

The build process will ensure that all the interpreters are installed and configured for your environment.

1.3 CUTS Analysis Toolset

The CUTS analysis tools allow developers to view and analyze system performance metrics. The analysis tools are developed using the following key technologies:

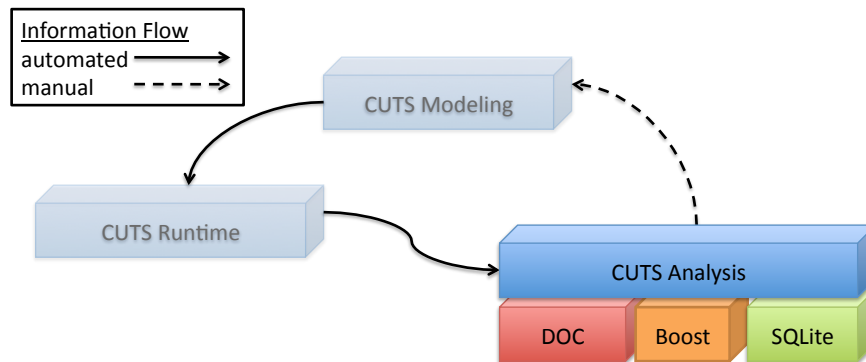


Figure 1.3. Building blocks for the CUTS analysis tools

As shown in Figure 1.3, the same technologies used to develop the runtime toolset (see Figure 1.1) are used to develop the analysis toolset. Because the same technologies are used in the analysis toolset, the core projects in analysis toolset is also built while building the runtime toolset.

Chapter 2

Quick Start Tutorial

This chapter provides a quick start tutorial for using CUTS. After reading this chapter, you will have a basic understanding of how to:

1. Model application behavior and workload.
2. Generate a test system model the constructed model.
3. Execute the system in your target environment.
4. Collect and analyze test results (COMING SOON).

In this tutorial, you will be measure and analyze the service time a simple server component. This tutorial assumes you have basic understanding of the Generic Modeling Environment (GME), CoSMIC/PICML, and the CORBA Component Model (CCM). This tutorial targets the Component Integrated ACE ORB (CIAO) middleware¹. Although this tutorial targets CIAO, the experience gained can be applied to other architectures that CUTS supports.

2.1 Modeling Behavior and Workload

Using GME, open the following file:

```
$ (CUTS_ROOT) /examples/MDE/GME/PICML/GettingStarted/GettingStarted.xme
```

This model contains the structure of the client/server application, and is usually the starting point for using CUTS within PICML. Currently, application behavior and workload is modeled in the Behavior aspect of a component's interface definition in PICML. The components in the model located at:

¹CIAO is an open-source implementation of the CCM, and is freely available for download at the following location: www.dre.vanderbilt.edu/CIAO.

already contain model elements for its input/output ports. What remains is associating the correct behavior and workload with these input/output ports for emulation purposes. First, let's add behavior and workload to the client component by assuming the client component should periodically send an event to the server. Please complete the following steps (elements will be automatically be generated as well):



1. Add a `PeriodicEvent` model element to the active model and set its name to `periodicPing` and its `Hertz` attribute to 10 (*i.e.*, 10 events/sec).
2. Insert an `InputAction` element, change its name to `periodicPing`, and connect the `PeriodicEvent` to the `InputAction`.
3. Insert an `OutputAction`, connect it with the `State`.
4. Connect the final `State` with the originating `InputAction`, *i.e.*, `periodicPing`, to signify end of the behavior.

Figure 2.1 illustrates the complete behavior model of the client in PICML, which is implemented in the previous steps.

Since you will need to collect performance metrics from the server component, its behavior/workload has already been completed. You, however, will learn how to add the necessary elements to the model to collect performance metrics in later chapters.

2.2 Generating Test System Implementation

After modeling the behavior and workload, the next step is to generate source code from the model. This will enable emulation of the test system on its target architecture. To generate source code from the model, launch the CUTS interpreter and execute the following steps (also illustrated in Figure 2.2):

1. Select `Generate component implementation` radio button;
2. Enter the target output directory;
3. Select `Component Integrated ACE ORB (CIAO)` in the listbox;
4. Click the OK button.

Once the CUTS interpreter finishes generating source code, the IDL and CIDL files need to be generated in order to compile the source code. These files are created by the IDL Generator and CIDL Generator interpreters, respectively, provided with CoSMIC ². Finally, there will be a `GettingStarted.mwc` file (its name may be different) located in the output directory selected for source code generation. This is a Makefile, Project and Workspace Creator (MPC) workspace file that contains all the necessary information to successfully compile the generated source code. Use `GettingStarted.mwc` to generate the appropriate workspace and then compile it.

2.3 Executing in Target Environment

One design goal of CUTS is to (re)use the same infrastructure used in the target environment. The `GettingStarted.xme` example uses the Deployment And Configuration Engine (DAnCE), which is included with CIAO's standard distribution, to deploy the test system. To deploy the example, use the following steps ³:

²Please ensure to generate the IDL and CIDL files in the same directory as the source code previously generated from the model.

³CUTS has a distributed testing framework that simplifies many of the complexities associated with running distributed system tests. We will cover the CUTS testing framework in later chapters

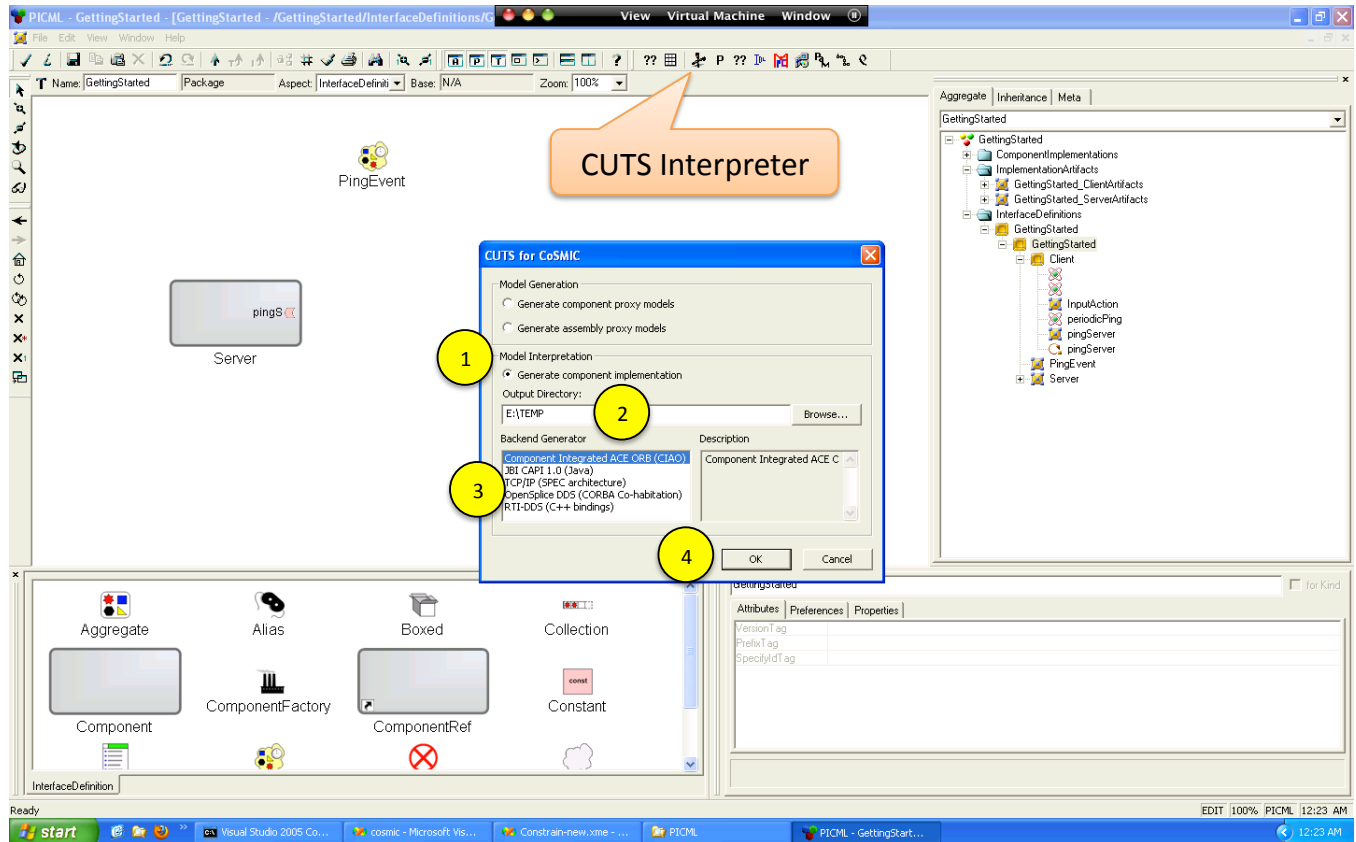


Figure 2.2. Walkthrough of the steps for generating source code from CBML/WML models in CUTS.

1. In the directory where you generated the source code for the example, create a directory named `./descriptors`.
2. Generate the deployment plan descriptors from the `GettingStarted` model into the `./descriptors` directory. Figure 2.3 shows the location of the deployment plan interpreter in CoSMIC.
3. Copy the files from `$CUTS_ROOT/examples/MDE/GME/PICML/GettingStarted` in the source distribution to the `./descriptors` directory.
4. Execute `run_test.pl` to deploy the system using DAnCE.

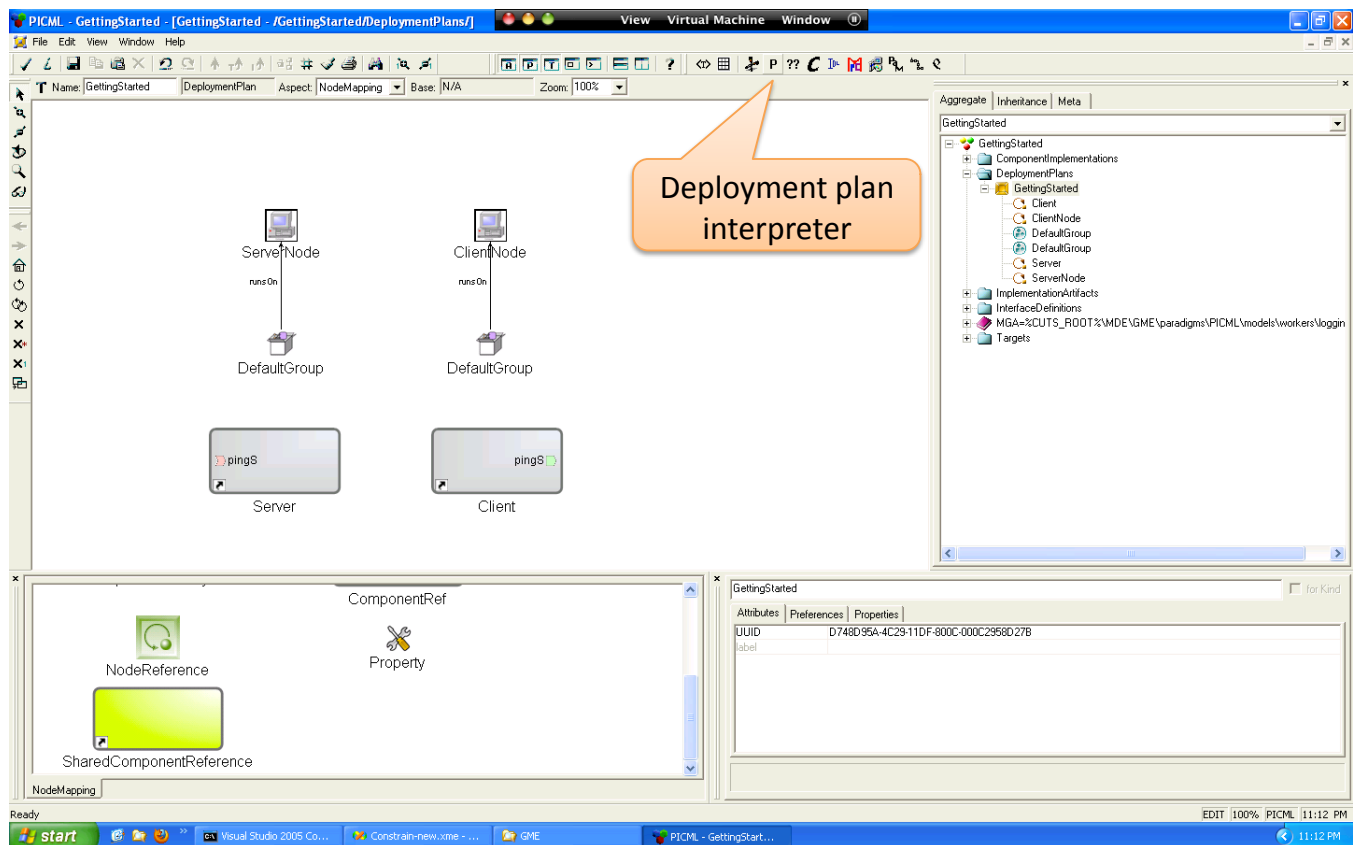


Figure 2.3. Location of the deployment plan interpreter in CoSMIC.

After the application is deployed, the client will send 10 events/second to the server. After 10 seconds, the script will teardown the application.

2.4 Analyzing Test Results

This section is coming soon.

Part II

CUTS Modeling Toolset

Part III

CUTS Runtime Toolset

Chapter 3

CUTS Node Management Facilities

This chapter discusses the CUTS node manager facilities. These facilities are designed to simplify management of many processes that must execute on many different nodes in a large-scale distributed testing environment. Furthermore, these facilities simplify management of many different execution environments, which is hard to do using traditional shell scripting facilities, such as Bash, C-Shell, and Batch files on Windows.

3.1 Overview

Distributed testing and execution environments, such as network testbeds, consist of many nodes. These nodes are responsible for executing many different processes that must support the overall goals of the execution environment. For example, nodes in the execution environment may host process that manage software components, processes that synchronize clocks, and process that collect software/hardware instrumentation data.

When dealing with many process that must execute on a given node, the traditional approach is to create a script that launches all processes for that given node. For example, in *Unix environments, Bash scripts can be used to launch a set of processes. The advantage of using such scripts is that is provides a *repeatable* technique for ensuring that all process for a given node are launched, and done so in the correct order. Moreover, the script is persistent and can be easily updated to either add or remove processes as needed.

Although such scripts suffice when working in a distributed testing environment, its approach as more complexity to the distributed testing process. First, the scripts are only repeatable after *manually* invoking the script. This means that distributed testers much manually connect to each host in the environment and (re)start the script. Even if its possible to remotely connect to the host via automation to (re)start the script, it is still *hard* to know what processes are associated with a given script.

Secondly, such scripts do not easily support injection/removal of processes on-the-fly. It is assumed the script is used to configure the host. At that point, all processes injected/removed from the host are assumed to be part of all the processes on that machine. Lastly, it is *hard* to switch between different execution environments where an execution environment is set of environment variables and a set of processes.

To address the limitations of traditional approaches to executing and managing a distributed testing environment, CUTS provides node management facilities. The facilities provide methods for managing many different execution environments, which can easily be swapped between. The facilities also provide tools for remotely controlling nodes in the execution environment.

The remainder of the chapter therefore presents the basics for using the CUTS node management facilities in a production-like environment. In particular, the following information is covered in this chapter:

- **CUTS Node Daemon** - The CUTS Node Daemon for managing many different execution environments for a node. Details of the CUTS Node Daemon are presented in Section 3.2.
- **CUTS Node Daemon Client** - The CUTS Node Daemon Client is responsible for remotely controlling the CUTS Node Daemon. Details of the CUTS Node Daemon Client are provided in Section 3.3.

3.2 The Node Daemon

The CUTS Node Daemon is responsible for managing different execution environments for a node. There is typically one CUTS Node Daemon executing on a single node. There can be cases where you have more than one CUTS Node Daemon executing on the same machine, but this defeats the purpose of the daemon. Regardless of the situation, what is discussed in the remainder of this chapter can be used to execute one or more CUTS Node Daemon processes on a single host.

3.2.1 Configuring the Node Daemon

The CUTS Node Daemon is designed to manage different execution environments. Before it can do this, however, it must be properly configured. The configuration for the CUTS Node Daemon is similar to a shell script, except that it is specified in XML format. Listing 3.1 illustrates an example configuration for the CUTS Node Daemon.

```

1 <?xml version="1.0" encoding="utf-8" standalone="no" ?>
2 <cuts:node xmlns="http://cuts.cs.iupui.edu"
3           xmlns:cuts="http://cuts.cs.iupui.edu"
4           xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
5           xsi:schemaLocation="http://cuts.cs.iupui.edu cutsnode.xsd">
```

```

6
7 <environment id="default" inherit="true" active="true">
8   <!-- environment variables for this environment -->
9   <variables>
10    <variable name="NameServiceIOR"
11           value="corbaloc:iiop:localhost:60003/NameService" />
12   </variables>
13
14   <startup>
15    <!-- NamingService -->
16    <process id="dance.naming.service" waitforcompletion="false">
17      <executable>${TAO_ROOT}/orbsvcs/Naming_Service/Naming_Service</executable>
18      <arguments>-m 0 -ORBEndpoint iiop://localhost:60003 -o ns.iior</arguments>
19    </process>
20
21    <!-- List of DAnCE NodeManager executions -->
22    <process id="dance.nodemanager.pingnode" waitforcompletion="false">
23      <executable>${DANCE_ROOT}/bin/dance_node_manager</executable>
24      <arguments>-ORBEndpoint iiop://localhost:30000
25                -s ${DANCE_ROOT}/bin/dance_locality_manager
26                -n PingNode=PingNode.iior -t 30
27                --instance-nc corbaloc:rir:/NameService</arguments>
28      <workingdirectory>../lib</workingdirectory>
29    </process>
30
31    <!-- The one and only execution manager -->
32    <process id="dance.executionmanager" waitforcompletion="false">
33      <executable>${DANCE_ROOT}/bin/dance_execution_manager</executable>
34      <arguments>-eEM.iior --node-map PingPong.nodemap</arguments>
35    </process>
36   </startup>
37 </environment>
38 </cuts:node>

```

Listing 3.1. An example configuration file for the CUTS Node Daemon.

Without going into the full details of the configuration in Listing 3.1, the CUTS Node Configuration contains three critical sections:

- **Environment** – The environment section, denoted by the `<environment>` tag, defines an execution environment. There can be more than one execution environment in a configuration. As shown in Listing 3.1, there is only one execution environment named `default`.

The remaining details of the execution environment is discussed in Section 3.2.4.

- **Variables** – The variables section, denoted by the `<variable>` tag, defines environment variables for the parent execution environment. All processes in the execution environment can use these environment variables via the `${}` tag (see line 17). As shown in Listing 3.1, the default execution environment contains one environment variable named `NameServiceIOR`. This environment variable is not explicitly used in the configuration. Instead, the processes in this configuration expect for the variable to be defined as part of the default environment variables. The remaining defaults of the variables section is discussed in Section 3.2.4.
- **Startup/Shutdown Processes** – The startup/shutdown section, denoted by `<startup>` and `<shutdown>`, respectively, defines the processes to execute when an execution environment starts up or shuts down. As shown in Listing 3.1, the configuration defines only a set of startup processes. The remaining details of the processes section is discussed in Section 3.2.4.

3.2.2 Running the Node Daemon

Assuming the CUTS runtime architecture has been built and installed correctly, the CUTS Node Daemon is installed at the following location:

```
%> $CUTS_ROOT/bin/cutsnode_d
```

To see a complete list of command-line options, use the following command:

```
%> $CUTS_ROOT/bin/cutsnode_d --help
```

Assuming the configuration discussed above is defined in a file named *cutsnode.config*, you run the CUTS Node Daemon with the configuration via the following command:

```
%> $CUTS_ROOT/bin/cutsnode_d -c cutsnode.config
-ORBEndpoint iiop://node.hydra.iupui.edu:20000
```

3.2.3 Configuring the Node Daemon Server

When you run the CUTS Node Daemon, the server is configured with its default values. In many cases, you may want to change the server's configuration. For example, you may want to bind the server to localhost instead of its hostname, or you may want to change the port on which it is listening.

You can configure the server using the `--server-args` command-line option. This should be a quoted command-line option because there can many different arguments passed to the server. The following table lists the different command-line options:

| Name | Required | Description |
|--------------------------|----------|--|
| --register-with-iortable | No | Name used to expose server via the IORTable |
| --register-with-ns | No | Name used to register server with naming service |

In addition to these options, you can pass any option supported by the ORB, such as `-ORBEndpoint`. The following command-line illustrates how to change the listening address of the CUTS Node Daemon to `node.hydra.iupui.edu:20000` and expose it via the IORTable as `CUTS/NodeDaemon`.

```
%> $CUTS_ROOT/bin/cutsnode
--server-args="-ORBEndpoint iiop://node.hydra.iupui.edu:20000
--register-with-iortable=CUTS/NodeDaemon"
```

3.2.4 Node Daemon Schema Definition

The section discusses the details of the CUTS Node Daemon configuration based on its underlying XML Schema Definition.

<cuts:node>

This is the main tag for the configuration document. Its XML namespace must have the `http://cuts.cs.iupui.edu` definition. If it does not have this definition, then the configuration will not load.

<environment>

This tag defines the different execution environments. As previously stated, there can be more than one execution environment in a configuration file. The following is a list of attributes supported by this tag.

| Name | Type | Default Value | Required | Description |
|---------|---------|---------------|----------|--|
| id | String | | Yes | Unique id for the execution environment |
| inherit | Boolean | false | No | Inherit the system's environment variables |
| active | Boolean | false | No | This is the active execution environment |

The following is a list of child tags: `<variables>`, `<startup>`, `<shutdown>`.

<variables>

This is the container tag for the environment variables in execution environment. This tag does not contain any attributes. The following is a list of child tags: `<variable>`.

<variable>

This tag is used to defined environment variables for an execution environment. The following is a list of attributes supported by this tag.

| Name | Type | Default Value | Required | Description |
|-------|--------|---------------|----------|-----------------------------------|
| name | String | | Yes | Name of the environment variable |
| value | String | | Yes | Value of the environment variable |

This tag does not contain any child tags.

<startup>

This tag is a container for a list of processes that are to executed with an environment starts up. The order in which the processes are listed is their startup order. The following is a list of child tags: <process>.

<shutdown>

This tag is a container for a list of processes that are to executed with an environment is shutdown. The order in which the processes are listed is their shutdown order. The following is a list of child tags: <process>.

<process>

This tag defines an executable process for the execution environment. The following is a list of attributes supported by this tag.

| Name | Type | Default Value | Required | Description |
|-------------------|---------|---------------|----------|---|
| id | String | | Yes | Unique id for the process |
| delay | Float | 0.0 | No | Seconds to wait before executing process |
| waitforcompletion | Boolean | false | No | Wait for process to complete before progressing |

The following is a list of child tags: <executable>, <arguments>, <workingdirectory>.

<executable>

This tag defines the complete/relative path of the executable process. Environment variables can be used within this tag. This tag does not contain any child tags.

<arguments>

This tag defines the command-line for the executable process. Environment variables can be used within this tag. This tag does not contain any child tags.

<workingdirectory>

This tag defines the working directory for the executable process (*i.e.*, the directory where executable process will be spawned). Environment variables can be used within this tag. This tag does not contain any child tags.

3.3 The Node Daemon Client

The purpose of the CUTS Node Daemon Client is to remotely control different CUTS Node Daemon application. This prevents users from having to manually login to different nodes when there is need to alter the execution environment.

3.3.1 Running the Node Daemon Client

Assuming the CUTS runtime architecture has been built and installed correctly, the CUTS Node Daemon Client is installed at the following location:

```
%> \ $CUTS\_ROOT/bin/cutsnode
```

To see a complete list of command-line options, use the following command:

```
%> \ $CUTS\_ROOT/bin/cutsnode --help
```

Assuming the CUTS Node Daemon is running at the address `iiop://node.hydra.iupui.edu:20000` and registered with the IORTable as `CUTS/NodeDaemon` as discussed above, then you can connect and control it via the CUTS Node Daemon Client using the following command:

```
%> $CUTS_ROOT/bin/cutsnode --reset -ORBInitRef
NodeDaemon=corbaloc:iiop:node.hydra.iupui.edu:20000/CUTS/NodeDaemon
```

In this example, the `-ORBInitRef` is a `corbaloc:.` This, however, need not always be the case. It can be any StringIOR support by *The ACE ORB (TAO)*, such as a IOR file or naming service location.

Chapter 4

CUTS Testing Facilities

Chapter 5

CUTS Logging Facilities

This chapter discusses the CUTS logging facilities. These facilities are designed to simplify collection of system execution traces in a distributed environment. Moreover, the facilities simplify partitioning system execution traces with different tests. This therefore makes it possible to execute multiple tests in the same environment while preserving data integrity. Finally, the logging facilities discussed in this chapter are extensible via interceptors, and are highly configurable so they adapt to different application domains.

5.1 Overview

Distributed systems consist of many software components executing on many hosts that are connected via a network. When validating distributed system function and quality-of-service (QoS) properties, *e.g.*, end-to-end response time, scalability, and throughput, it is necessary to collect data about the systems behavior in the target environment. Such behaviors could be the events executed of the execution lifetime of the system, the state of the system at different points in time, or data points needed to calculate the end-to-end response of an event.

When dealing with data collection in a distributed environment, such as those that host distributed systems, data is collected and analyzed either *offline* or *online*. In offline collection and analysis, data from each host is written to local persistent storage, such as a file, while the system is executing in its target environment. After the system is shutdown, the collected data in local storage on each host is combined and analyzed. The advantage of this approach is network traffic is kept to a minimum since collected data is not transmitted over the network until after the system is shutdown. The disadvantage of this approach is collected data is not processed until the system is shutdown. This can pose a problem for systems with a long execution lifetime, *e.g.*, ultra-large-scale systems that must run 24x7, or when trying to monitor/analyze a system in real-time using collected data.

In online collection and analysis, data is collected and transmitted via network to a host outside the execution environment. The advantage of this approach is that it allows analysis of metrics data an environment that does not use the systems resources. This also helps not to skew the results while the system is running. The disadvantage of online collection and analysis is the difficulty of devising a strategy for efficiently collecting data in a distributed environment and submitting it to a central location without negatively impacting the excuting systems QoS—especially if the system generates heavy network traffic.




Figure 5.1. Overview of the CUTS logging facilities.

Because online collection and analysis is a more promising approach to data collection and analysis, CUTS logging facilities implement online data collection. As illustrated in Figure 5.1, there is a target execution environment and a runtime environment. The runtime environment is where the system is executing and generating data for collection. Each host in the runtime environment has a logging client (see Section 5.3). The logging client is responsible for collecting data from individual software/hardware components on its host. After collecting data on the host, the logging client periodically submits collected data to a logging server (see Section 5.2), which is executing outside of the runtime environment.

The remainder of the chapter therefore presents the basics for using the CUTS logging facilities in a production-like environment. In particular, the following information is covered in this chapter:

- **CUTS Logging Server** - The CUTS Logging Server is responsible for collecting data from individual CUTS Logging Clients executing in the runtime environment. Details of the CUTS Logging Server are presented in Section 5.2.

- **CUTS Logging Client** - The CUTS Logging Client is responsible for collecting data from individual client loggers and submitting it to the CUTS Logging Server. Details of the CUTS Logging Client are provided in Section 5.3.
- **Client Logger** - The client logger is responsible for collecting actual data from the runtime environment and submitting it to the CUTS Logging Client executing on its host. Details of the client logger are provided in Section 5.4.

5.2 The Logging Server

The logging server is responsible for collecting data from individual logging clients in the runtime environment (see Section 5.3). As shown in Figure 5.1, there is typically one instance of a logging server running in the target execution environment. There can also be cases when more than one instance of a logging server is desired, such as using multiple logging servers to balance the workload due to the amount of data collected by individual logging clients in the runtime environment. This user's manual, however, assumes that only a single instance of a logging server is used in the runtime environment.

5.2.1 Running the Logging Server

Assuming the CUTS runtime architecture has been built and installed correctly, the CUTS Logging Server is installed at the following location:

```
%> $CUTS_ROOT/bin/cuts-logging-server
```

To see a complete list of command-line options, use the following command:

```
%> $CUTS_ROOT/bin/cuts-logging-server --help
```

5.2.2 Configuring the Logging Server

There are only a handful of command-line options for configuring the logging server. The most important of the available command-line options is `--register-with-iortable=NAME`, where NAME is an user-defined name, such as `LoggingServer`. When this option is combined with the `-ORBEndpoint ENDPOINT CORBA` command-line option, an external reference to the logging server is exposed. This is what allows the CUTS Logging Client (see Section 5.3) to connect to the logging server.

The following is an example of exposing the logging server named `LoggingServer` on port 20000 for logging clients to connect:

```
%> $CUTS_ROOT/bin/cuts-logging-server -ORBEndpoint \
    iiop://`hostname`:20000 --register-with-iortable=LoggingServer
```

As shown in the example above, the logging server is listening using port 20000 on its host ¹, and has the name `LoggingServer`. The logging client can therefore connect to the logging server using the following CORBA reference:

```
corbaloc:iiop:`hostname`:20000/LoggingServer
```

The next section therefore discusses how to configure the logging client to connect to the correct logging server in the runtime environment.

5.3 The Logging Client

The logging client is responsible for collecting data from individual client loggers in the runtime environment (see Section 5.4). As shown in Figure 5.1, each host in the runtime environment typically runs one instance of a logging client. Similar to the logging server, there can also be cases when more than one instance of a logging client is desired, such as using multiple logging clients to balance the workload due to the amount of data collected by individual client loggers in the runtime environment. This user's manual, however, assumes that only a single instance of a logging client running on each host in the runtime environment.

5.3.1 Running the Logging Client

Similar to the CUTS Logging Server, the CUTS Logging Client is installed at the following location:

```
%> $CUTS_ROOT/bin/cuts-logging-client
```

To see a complete list of command-line options, use the following command:

```
%> $CUTS_ROOT/bin/cuts-logging-client --help
```

5.3.2 Configuring the Logging Client

The CUTS Logging Client is configured similar to how the CUTS Logging Server is configured (see Section 5.2.2). For example, the `-ORBEndpoint ENDPOINT` defines the CORBA endpoint for the logging client. The logging client's endpoint is used by the client loggers (see Section 5.4) that need to connect to the CUTS Logging Client executing on its local host.

¹'hostname' is used because it is assumed that Unix command-line substitution is available on the host in the example. If such a feature is not available, it is possible to hardcode the hostname in place of 'hostname', e.g., `host.foo.bar.com`

The addition to configuration parameters similar to the CUTS Logging Server, to the logging client's configuration requires the location of the CUTS Logging Server. In Section 5.2.2, the logging server was given an endpoint that is used by the logging client when connecting to it. To specify the location of the CUTS logging server, use the `-ORBInitRef LoggingServer=LOCATION` CORBA command-line option, where `LOCATION` is the location of the logging server².

The following is an example of configuring the logging client to connect to the logging server discussed in Section 5.2.2 and listens to data from client loggers on port 20000:

```
%> $CUTS_ROOT/bin/cuts-logging-client -ORBEndpoint \
    iiop://`hostname`:20000 --register-with-iortable=LoggingClient \
    -ORBInitRef LoggingServer=corbaloc:iiop:server.foo.bar.com:20000/LoggingServer
```

The next section discusses how to use client loggers to submit data for collection and analysis to the logging clients executing on each host in the runtime environment.

5.4 The Client Loggers

The client loggers are responsible for collecting log messages from an application and submitting them to the logging client on its host machine. Since client loggers interact directly with the application, they are language dependent. Currently, CUTS supports client loggers for C++ and Java applications. In addition, CUTS support the following logging frameworks: ACE Logging Facilities and log4j. The remainder of the section discusses how directly integrate client loggers into an application and indirectly intergrate them using an pre-existing logging framework.

5.4.1 Direct Integration

Direct integration is the simplest method for integrating CUTS logging facilities into an existing application. It is also the most intrusive approach for integrating the logging facilities into an existing application. This is because you have to augment existing code to include the client logger(s). Only then can the application use to client logger to collect messages and submit them to the logging client executing on the localhost.

The following code illustrates integrating the client logger using the C++ version:

```
1 #include "cuts/utils/logging/client/logger/Client_Logger.h"
2
3 int main (int argc, char * argv []) {
4     try {
5         // connect client logger to logging client.
6         CUTS_Client_Logger logger;
```

²The location of the `LoggingServer` can be a `corbaloc`, `corbaname`, or `IOR`.


```

7      if (-1 == logger.connect (argv[1]))
8          return 1;
9
10     // Send a single, yet simple message.
11     logger.log (LM_DEBUG, "this is a simple message");
12
13     // Disconnect from the logging client.
14     logger.disconnect ();
15     return 0;
16 }
17 catch (const ::CORBA::Exception & ex) { }
18 catch (...) { }
19
20 return 1;
21 }

```

Listing 5.1. Example illustrating integration of the C++ client logger.

Likewise, this code illustrates the same concept using the Java version of the client logger:

```

1  import CUTS.client.logging.Logger;
2
3  public class TestApplication {
4      public static void main (String [] args) {
5          String loggingClientLoc = /* location of logging client */
6          Logger clientLogger = new Logger ();
7
8          // connect to the logging client
9          clientLogger.connect (loggingClientLoc);
10
11         // send a log message
12         clientLogger.log ("this is a simple message");
13
14         // disconnect from the logging client
15         clientLogger.disconnect ();
16     }
17 }

```

Listing 5.2. Example illustrating integration of the Java client logger.

In both the C++ and Java example, the `connect ()` method takes as its input parameter the location of the logging client. This location can be a CORBA IOR, corbaname, or corbaloc (*i.e.*, anything that is a valid parameter for the `string_to_object ()` method implemented on the

ORB). After a connection is made to the logging client, the client logger submits messages to the logging client using the `log()` method. When the client logger is done (or the application is ready to exit), it disconnects from the logging client using the `disconnect()` method. This is a critical step because it ensures the logging client does not get into an inconsistent state when trying to keep track of the number of active client loggers.

5.4.2 Indirect Integration

Indirect integration is a non-intrusive method for integrating the CUTS logging facilities with an existing application. This approach only work if the application is already using a logging framework, such as ACE Logging Facilities or log4j. Indirect integration works by intercepting log messages sent to the original logging framework and submitting them to the CUTS logging facilities (*i.e.*, the logging client).

Unlike direct integration, system developers do not need to update their existing code base to leverage indirect integration. Instead, they must update the existing logging frameworks configuration files to “install” the necessary inteceptors. The remainder of this section therefore presents details on using indirect integration with supported logging frameworks in CUTS.

log4j

Language. Java

How to integrate. Update `log4j.properties` (or similar file)

Illustrative example.

```
# define the loggers
log4j.rootCategory=ALL, Console, LoggingClient

# console appender
log4j.appender.Console=org.apache.log4j.ConsoleAppender
log4j.appender.Console.layout=org.apache.log4j.PatternLayout
log4j.appender.Console.layout.ConversionPattern=%-4r [%t] %-5p %c %x - %m%n

# CUTS appender
log4j.appender.LoggingClient=cuts.log4j.LoggingClientAppender
log4j.appender.LoggingClient.LoggerClient=corbaloc:iiop:localhost:20000/LoggingClient
```

ACE Logging Facilities

Language. C++

How to integrate. Update `svc.conf` (or similar file)

Illustrative example.

```
dynamic CUTS_ACE_Log_Interceptor Service_Object * \  
CUTS_ACE_Log_Interceptor::_make_CUTS_ACE_Log_Interceptor() active \  
"--client=corbaloc:iiop:localhost:20000/LoggingClient"
```

Part IV

CUTS Analysis Toolset

Part V

Appendix

Appendix A

Building and Installing Third-Party Libraries

The following appendix contains best practices for building and installing third-party libraries used by CUTS projects. It is not a requirement to install third-party libraries per the instructions in this appendix. If you are experiencing compilation problems with CUTS due to third-party libraries, however, the instructions in this appendix may be of assistance. To make help explain the installation process of the thirdparty libraries, it will be assumed that all third-party libraries will be installed at the following location: `/opt/thirdparty`^{1 2}.

A.1 Boost

The Boost libraries are required when building and installing CUTS. All artifacts for Boost, such as prebuilt versions and source files, can be downloaded from the following location: <http://www.boost.org>.

A.1.1 System Configuration

To build CUTS correctly, the Boost environment variables must be defined. Please define the following required environment variables:

```
%> export BOOST_ROOT= location of Boost
%> export BOOST_VERSION= name of subdirectery under $BOOST_ROOT/include, e.g.,
boost-1_34_1
```

¹This location is representative of a common location for installing third-party libraries, and may be different on your system.

²If you are installing CUTS from source on Windows, it is recommended that you place all environment variables in a BATCH file (`.bat`) so you do not pollute the system environment.


```
%> export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$BOOST_ROOT/lib3
%> export PATH=$PATH:$BOOST_ROOT/lib
```

In some cases, Boost libraries will have an optional configuration appended to their filename, such as `-mt-d`. To use a specific configuration of Boost, the `BOOST_CFG` environment variable must be defined. The following is an example of defining the Boost configuration environment variable:

```
%> export BOOST_CFG=-mt-d
```

The configuration of Boost is now complete. Once all the other third-party libraries are installed you will be able to build CUTS (see Chapter 1).

A.1.2 Installation

Non-Windows

If you are installing Boost on a non-Windows system, *e.g.*, Linux, Solaris, or MacOS X, then you can download the source code directly from its main website (<http://www.boost.org>), build it, and install it using standard procedures.

Windows

Installing Boost from sources on Windows is not a trivial process. Moreover, the prebuilt version of Boost available for download from the main website does not contain the correct configuration or directory layout. It is therefore recommended that you download and use the prebuilt version of Boost available for download from the CUTS website at the following location: <http://www.dre.vanderbilt.edu/CUTS/downloads/thirdparty/boost>. Please download the correct version of Boost that matches your version of Microsoft Visual C++ and extract it to a common location on disk.

A.2 Xerces-C

Xerces-C 3.x is required when building and installing CUTS and other third-party libraries. All artifacts for Xerces-C, such as prebuilt versions and source files, can be downloaded from the following location: <http://xerces.apache.org/xerces-c>.

A.2.1 System Configuration

To build CUTS correctly, the Xerces-C environment variables must be defined. Please define the following required environment variables:

³On Windows-based systems, please add library paths to `PATH`. On Mac OS-based systems, please add library paths to `DYLD_LIBRARY_PATH`.

```
%> export XERCESSROOT= location of Xerces-C
%> export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$XERCESSROOT/lib
```

The configuration of Xerces-C is now complete. Once all the other third-party libraries are installed you will be able to build CUTS (see Chapter 1).

A.2.2 Installation

Non-Windows

If you are installing Xerces-C on a non-Windows system, *e.g.*, Linux, Solaris, or MacOS X, then you can download the source code directly from its main website (<http://xerces.apache.org/xerces-c>), build it, and install it using standard procedures.

Windows

Installing Xerces-C from sources on Windows is not a trivial process. Moreover, the prebuilt version of Xerces-C available for download from the main website does not contain the correct configuration or directory layout. It is therefore recommended that you download and use the prebuilt version of Xerces-C available for download from the CUTS website at the following location: <http://www.dre.vanderbilt.edu/CUTS/downloads/thirdparty/xerces-c>. Please download the correct version of Xerces-C that matches your version of Microsoft Visual C++ and extract it to a common location on disk.

A.3 DOC Group Middleware

DOC Group Middleware is required when building and installing CUTS. Its source code can be downloaded from the following location: <http://www.dre.vanderbilt.edu/CIAO>.

A.3.1 System Configuration

Unlike Boost, you must manually configure your system before installing the DOC middleware. Once you have downloaded the tarball from the website above⁴, please define the following environment variables:

```
%> export ACE_ROOT= location of ACE
%> export TAO_ROOT= location of TAO
%> export CIAO_ROOT= location of CIAO
```

⁴It is recommended you download either the latest version or snapshot of trunk from the SVN repo. Otherwise, you will not have access to ADBC.

```
%> export ADBC_ROOT= location of ADBC
%> export DANCE_ROOT=$CIAO_ROOT/DAnCE

%> export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$ACE_ROOT/lib:$ADBC_ROOT/lib
%> export PATH=$PATH:$ACE_ROOT/bin:$CIAO_ROOT/bin:$DANCE_ROOT/bin
```

After defining the environment variables above, the next step is configuring the middleware for the target platform. This is accomplished by defining the `$ACE_ROOT/ace/config.h` header file, which will include the correct configuration for your platform. The included configuration file has the form `config-*`, where `*` is the platform of choice. The following is an example `config.h` file for building DOC middleware on Linux platforms:

```
// -*- C++ -*-

#ifdef _ACE_CONFIG_H_
#define _ACE_CONFIG_H_

#include ``config-linux.h``

#endif // !defined _ACE_CONFIG_H_
```

Non-Windows

Configuring the DOC middleware on non-Windows systems requires one more step. The configuration for the build engine also needs to be defined. Please use the following step to define the build engine's configuration:

```
%> cd $ACE_ROOT/include/makeinclude
%> ln -s platform*.GNU platform_macros.GNU
```

where `*` is the architecture for your platform of choice.

A.3.2 Installation

After you have configured your system, you are ready to build and install ACE and TAO. Please use the following steps to build and install the DOC middleware:

```
%> cd $CIAO_ROOT
%> $ACE_ROOT/bin/mwc.pl -type [build tool] CIAO-TAO-DAnCE.mwc
%> gmake
```

A.3.3 ACE DataBase Connector (ADBC) Framework

The ACE DataBase Connector (ADBC) Framework is a set of C++ wrappers that provide a common interface for using different database drivers. To install ADBC, first copy the file `default.features.tmpl` to `default.features`. Open the new file and then enable/disable the necessary features based on what wrappers you want to build.⁵ Finally, generate the workspace and build ADBC using the following commands:

```
%> cd $ADBC_ROOT
%> $ACE_ROOT/bin/mwc.pl -type [build tool] ADBC.mwc
%> gmake
```

A.4 XML Schema Compiler (XSC)

XSC is required when building and installing CUTS. All artifacts for XSC can be downloaded from the following location: <svn://svn.dre.vanderbilt.edu/XSC/trunk>. The remainder of this section discusses how to build and install XSC.

A.4.1 System Configuration

You must configure your system installing XSC. Once you have downloaded XSC from the website above please define the following environment variables:

```
%> svn co svn://svn.dre.vanderbilt.edu/XSC/trunk XSC
%> export XSC_ROOT= location of XSC
%> export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$XSC_ROOT/lib
%> export PATH=$PATH:$XSC_ROOT/bin
```

A.4.2 Installation

After you have configured your system for XSC, you are ready to build and install it. Please use the following steps to build and install XSC:

```
%> cd $XSC_ROOT
%> $ACE_ROOT/bin/mwc.pl -type [build tool] -features xsc=1,xerces3=1 XSC.mwc
```

This will generate the workspace for building and installing XSC. The name of the workspace depends on the build tool specified when generating the workspace. Finally, using your build tool of choice, you can build and install XSC.

⁵It is assumed that you have installed the development version for the database drivers that you plan to build.

A.5 Perl-Compatible Regular Expressions (PCRE)

PCRE is required when building and installing several projects in CUTS. You can download the source code for PCRE from the following location: <http://www.pcre.org>

A.5.1 System Configuration

To build CUTS correctly, the PCRE environment variables must be defined. Please define the following required environment variables:

```
%> export PCRE_ROOT= location of PCRE
%> export PCRE_VERSION= version of PCRE, blank if not applicable
%> export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$PCRE_ROOT/lib
```

The configuration of PCRE is now complete. Once all the other third-party libraries are installed you will be able to build CUTS (see Chapter 1).

A.5.2 Installation

Non-Windows

If you are installing PCRE on a non-Windows system, *e.g.*, Linux, Solaris, or MacOS X, then you can download the source code directly from its main website (<http://www.pcre.org>), build it, and install it using standard procedures.

Windows

Installing PCRE from sources on Windows is not a trivial process. It is therefore recommended that you download and use the prebuilt version of PCRE available for download from the CUTS website at the following location: <http://www.dre.vanderbilt.edu/CUTS/downloads/thirdparty/pcre>. Please download the correct version of PCRE that matches your version of Microsoft Visual C++ and extract it to a common location on disk.

A.6 SQLite

SQLite is required when building and installing several projects in CUTS. You can download the source code for SQLite from the following location: <http://www.sqlite.org>

A.6.1 System Configuration

To build CUTS correctly, the SQLite environment variables must be defined. Please define the following required environment variables:

```
%> export SQLITE_ROOT= location of SQLite
%> export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$SQLITE_ROOT/lib
%> export PATH=$PATH:$SQLITE_ROOT/bin
```

The configuration of SQLite is now complete. Once all the other third-party libraries are installed you will be able to build CUTS (see Chapter 1).

A.6.2 Installation

Non-Windows

If you are installing SQLite on a non-Windows system, *e.g.*, Linux, Solaris, or MacOS X, then you can download the source code directly from its main website (<http://www.sqlite.org>), build it, and install it using standard procedures.

Windows

Installing SQLite from sources on Windows is not a trivial process. It is therefore recommended that you download and use the prebuilt version of SQLite available for download from the CUTS website at the following location: <http://www.dre.vanderbilt.edu/CUTS/downloads/thirdparty/sqlite>. Please download the correct version of SQLite that matches your version of Microsoft Visual C++ and extract it to a common location on disk.

A.7 RTI-DDS

RTI-DDS is required if you are building CHAOS. You can obtain RTI-DDS from the following location: <http://www.rti.com>

A.7.1 System Configuration

In order to use RTI-DDS with CUTS, the environment must be configured a certain way. Otherwise, there is great chance that CUTS emulation code that uses RTI-DDS will not compile correctly. Please define the following environment variables:

```
%> export NDDSHOME= location of RTI-DDS
%> export NDDSARCHITECTURE= target arch (see $NDDSHOME/lib)
%> export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$NDDSHOME/lib/$NDDSARCHITECTURE
```

A.7.2 Installation

RTI-DDS comes prebuilt. There are no installation procedures once you unzip the obtained tarball.

A.8 OpenSplice DDS

OpenSplice DDS is required if you are building CHAOS. You can obtain OpenSplice DDS from the following location: <http://www.opensplice.org>. Please make sure to download the binary version.

A.8.1 System Configuration

In order to use OpenSplice DDS with CUTS, the environment must be configured a certain way. Otherwise, there is great chance that CUTS emulation code that uses OpenSplice DDS will not compile correctly. Please define the following environment variables:

```
%> export OSPL_HOME= location of OpenSplice HDE/target
%> export OSPL_TARGET= target arch (see $OSPL_HOME)
%> export OSPL_TEMP_PATH=$OSPL_HOME/etc/idlpp
%> export OSPL_URI="file://$OSPL_HOME/etc/config/ospl.xml"

%> export CLASSPATH=$OSPL_HOME/jar/dcpssaj.jar;$CLASSPATH

%> export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$OSPL_HOME/lib
```

A.8.2 Installation

OpenSplice DDS is built on top of The ACE ORB (TAO), which is also the case for CUTS. In most cases, however, the version of TAO in either OpenSplice DDS and CUTS is not the same version. Likewise, OpenSplice DDS and CUTS may not be built in such a way that they support different versions in the same process space, which is possible given the correct configuration. Because of this, it is best to rebuild only the custom library of OpenSplice DDS so that both OpenSplice DDS and CUTS can operate in the same address space.

To rebuild the custom library in OpenSplice DDS, first determine the version of TAO installed on your machine. The version number is found in the following file: `$TAO_ROOT/VERSION`. Once you have identified the version number, create the following directories in OpenSplice DDS install directory:

```
%> cp $OSPL_HOME/etc/idlpp/CCPP/[current version DDS_ACE_TAO]
$OSPL_HOME/etc/idlpp/CCPP/DDS_ACE_TAO_x-y-z
%> cp $OSPL_HOME/include/dcps/C++/CCPP/[current version DDS_ACE_TAO]
$OSPL_HOME/include/dcps/C++/CCPP/DDS_ACE_TAO__x-y-z
```

where `x_y_z` is the correct version of TAO installed on your machine. Please make sure to use underscores (`_`) and not dots (`.`) in the version number. Finally, define the following environment variable so the OpenSplice DDS architecture knows what version of the ORB to use:

```
%> export SPLICE_ORBE=DDS_ACE_TAO_x.y.z
```

where `x_y_z` is the correct version of TAO installed on your machine as done above.

The final step in the process is rebuilding the custom library is actually building it. This can be accomplished using the provided MPC project file in CUTS and executing the following steps:

```
%> cd $CUTS_ROOT/contrib/OpenSplice
%> $ACE_ROOT/bin/mwc.pl -type [build type]
%> #build the generated solution
```

The build process will install the custom library in `$OSPL_HOME/lib`.

Appendix B

Autobuilds of CUTS

This section contains detailed instructions for hosting an autobuild of CUTS. The autobuild is designed to both build and test core functionality of CUTS so errors can be located in a timely manner. Moreover, hosting an autobuild is designed to be an easy process. Please use these instructions as a starting point. If you need to change any part of the process, please feel free to do so. We ask, however, that you please do not commit your changes back to the repository.

B.1 Hosting a Windows Build

Windows-based builds of CUTS are conducted using CruiseControl (cruisecontrol.sourceforge.net), which is a Java-based continuous integration environment. The scripts for hosting a Windows build are already defined for CUTS. You can find them at the following location in the CUTS directory structure:

```
%> $CUTS_ROOT/integration
```

Once CruiseControl is installed in your environment, you can include CUTS in the build cycle in two main steps. First, you must include the CUTS project in CruiseControl's main configuration file (*i.e.*, `config.xml`). The following line is an example of including the Visual Studio 2005 build for CUTS as part of the build cycle.

```
<include.projects  
  file="E:/proj/vc8/CUTS/integration/cruisecontrol/cuts-vc8.config.xml" />
```

Configuring the build environment. After including the desired projects in CruiseControl's main configuration file, the next (and final) step is to configure the environment. This is done by first defining the directory structure for the build environment. Each build type (*e.g.*, Visual Studio 2005 and Visual Studio 2008) must have its own sandbox in the target build environment. The sandbox is where all projects for that build type must should reside when building CUTS

via CruiseControl. The following listing shows the directory structure for hosting both a Visual Studio 2005 and Visual Studio 2008 build on the same machine where the sandboxes are located in `E:/proj`.

```
/vc80
  /CUTS
  /ACE_TAO_CIAO
  /XSC
  ...
/vc90
  /CUTS
  /ACE_TAO_CIAO
  /XSC
  ...
```

Once you have established the sandbox locations, you must inform the build environment of its location. This can be done by setting one or more of the following environment variables. Using

Table B.1. Environment variables that capture the location of a sandbox.

| Variable Name | Description |
|---------------|---|
| VC80SANDBOX | Defines the sandbox for Visual Studio 2005 builds |
| VC90SANDBOX | Defines the sandbox for Visual Studio 2008 builds |

the example directory structure above, the Visual Studio 2005 would be defined as follows:

```
set VC80SANDBOX=E:/proj/vc8
```

Notice, this environment variable must be defined in a persistent location, such as the user/system environment variables.

Lastly, you must define the appropriate configuration environment variable. This is necessary because each build environment may install the required third-party libraries (see Appendix A) in different locations. The configuration environment variable(s) therefore points to a batch (`.bat`) file that defines all the environment variables need to build CUTS, such as all the environment variables defined in Appendix A. The following table is the list of configuration environment variables based on the build type:

Table B.2. Environment variables that define the build configuration.

| Variable Name | Description |
|---------------|---|
| VC80BUILDCFG | Defines build configuration for Visual Studio 2005 builds |
| VC90BUILDCFG | Defines build configuration for Visual Studio 2008 builds |

For example, if Visual Studio 2005 environment variables are defined in `E:/proj/vc9/setenv.bat`, then `VC80BUILDCFG` would be defined as follows:

```
set VC80BUILDCFG=E:/proj/vc8/setenv.bat
```

Notice, this environment variable must be defined in a persistent location, such as the user/system environment variables.

Congratulations, you are now ready to build Windows-based version of CUTS using CruiseControl!